Unified Status + Plan (One View)

1) Current status (as of your latest logs)
- Bot is running and scanning continuously.
- Last traded decision in the current run was around 2026-02-23 05:01 AM EST from traded-decisions-2026-02-23.json.
- Since then, evaluations continue, but most end in NO-TRADE due Kelly returning no position.

2) Why no trades are being placed now
- The primary blocker is the uncertainty gate:
  - uncertainty_haircut_max = "0.5" in config.toml.
  - Kelly hard-rejects when model_uncertainty > max_uncertainty in mod.rs.
- Current per-asset uncertainty values are high:
  - BTC 0.7007
  - ETH 0.6567
  - XRP 0.5357
  - SOL 0.9192

3) What per-asset uncertainty is
- It is a calibration reliability score per coin.
- Computed from rolling Brier score:
  - model_uncertainty = min(1.0, 2 * brier_score) in calibration.rs.
- Lower is better (more calibrated). Higher means model probabilities are less trustworthy.

4) Is the base model accurate enough right now?
- Infrastructure: yes (scanner, pipeline, settlement, P&L, logs are functioning).
- Predictive quality: not yet proven strong.
- With current uncertainty levels and near coin-flip outcomes, this is a good operational baseline, not yet a high-confidence trading model.

5) Stage-2 (fine-tuning) should look like this
- Persist calibration history across restarts (currently it resets each run).
- Keep per-asset calibrators (BTC/ETH/XRP/SOL) and update only on settled official outcomes.
- Add warmup logic:
  - Do not hard-block on uncertainty until minimum settled samples per asset.
- Use calibrated probabilities for entry and Kelly sizing.
- Keep sampled no-trade logs for counterfactual learning (important for improving filters).
- Validate with walk-forward time splits only (no random splits).
- Promote config/model only if out-of-sample EV and drawdown improve.

6) Immediate practical changes (to unblock testing now)
- Raise uncertainty_haircut_max from 0.5 to a temporary higher value (for example 0.85-0.95) in config.toml.
- Keep conservative sizing (kelly_fraction low) while calibrating.
- Track per-asset Brier/log-loss trend and win rate on closed trades only.

7) Best next implementation order
- Calibration persistence + load-on-start.
- Warmup-aware Kelly gate.

- Calibrator layer per asset.
- Walk-forward evaluation harness using your JSON logs.
- Policy tuning (confidence/spread/liquidity/time filters) from data, not guesswork.

Research references
- scikit-learn Calibration
- Guo et al. (2017) On Calibration
- TimeSeriesSplit
- Probability of Backtest Overfitting
- Online learning with delayed feedback